

# A Survey on SGX Enclave Privileged Side-Channel Attacks

Amin Fallahi

*Ph.D. Student*

---

Syracuse University

Department of Electrical Engineering and Computer Science

MAY 14, 2018

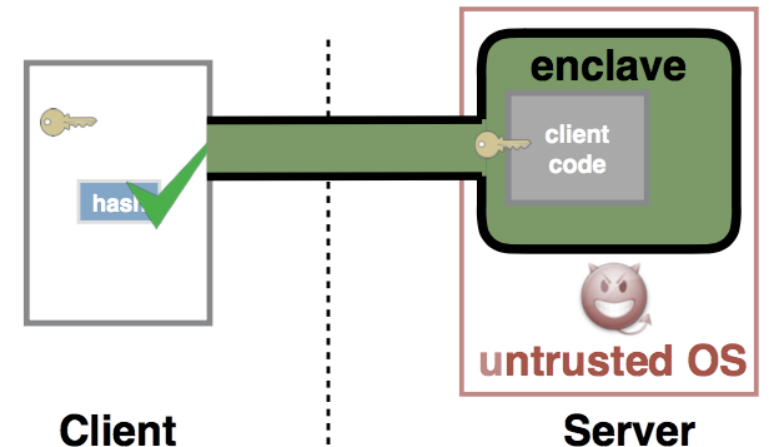
# Contents

---

- Introduction
- Attacks
- Defenses
- Analysis
- Conclusion

# Introduction

- Computing environment
  - Shared resources executing user programs
- Security inside the cloud
  - A hardware solution needed
- Intel SGX
  - Secure user code and data in the public cloud
  - Protecting data inside the enclave
- Privileged Operating System
  - Kernel space, hardware control



# Introduction

---

## *Intel SGX*

- 17 new instructions
- Enclave created by OS
  - Contiguous physical memory
  - Base address randomly assigned from EPC
- Hardware level memory encryption engine
- Virtual to physical address translation done by OS
- Interrupt  $\Rightarrow$  context switch – AEX
- ECALLs & OCALLs
- TPM: Trusted Platform Module
- TXT: Trusted Execution Technology

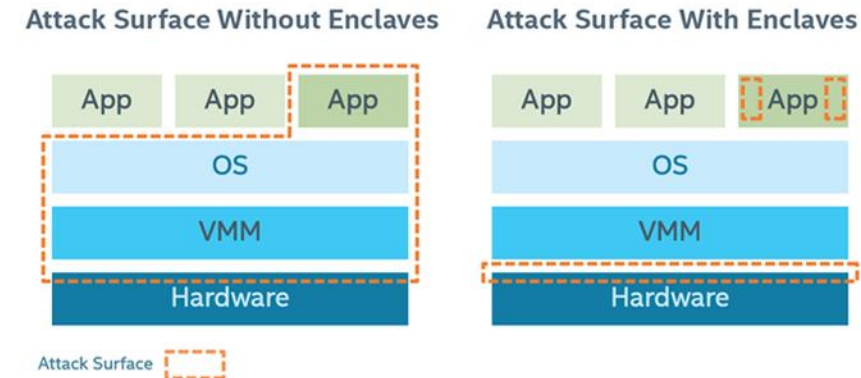
# Introduction

- Side-Channels

- Attacks by analyzing system behavior
- Cache attacks: Monitoring cache accesses
- Timing attacks: Measuring time between computations
- Page-fault attack: Monitoring page-faults and analyzing the pattern

- Attack Model

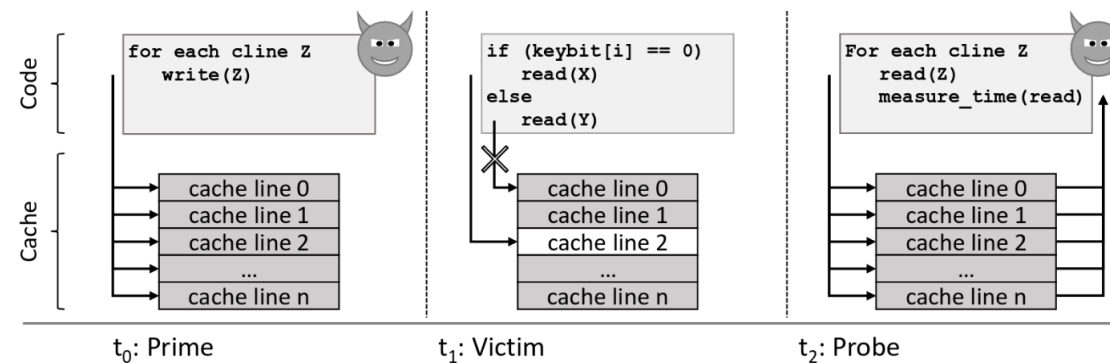
- Enclave protected by SGX
  - No stranger has the key
- Operating System can manage task queues, interrupts, and exceptions
- Viewing data transfer between memories and caches
- Program pinned to one core, no excessive interrupts, isolated



# Attacks

## Cache-Based Attacks

- Intentional Cache-Misses  $\Rightarrow$  Make program access memory
- Time, trace, access patterns
- Prime+Probe
  - Prime: attacker executes conflicting cache lines  $\Rightarrow$  cache miss
  - Probe: executing all the cache lines and measuring execution time



# Attacks

## Cache-Based Attacks

- RSA key extraction
  - Modular exponentiation
  - Large  $e$ ,  $d$ , and  $n$  such that  $(m^e)^d \equiv m \pmod{n}$
  - Knowing  $m$ ,  $e$ , and  $n$  it is extremely difficult to find  $d$
- Square and multiply
  - Typical algorithm  $\Rightarrow$  find the exponent
  - Secret dependent memory accesses
- Challenges
  - High precision timer
  - Eviction set generation


---

### ALGORITHM 3: Timestamp counter [23]

---

```
mov &counter, %rcx
l: inc %rax
mov %rax, (%rcx)
jmp l
```

---



---

### ALGORITHM 2: Square and multiply exponentiation [23]

---

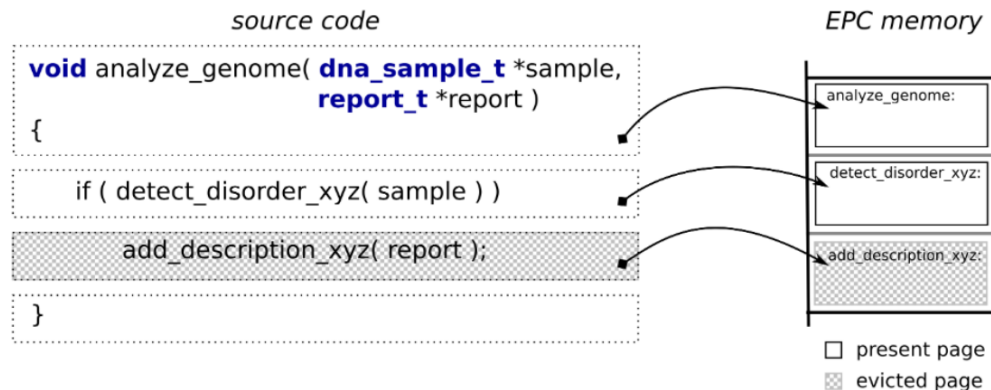
```
input: base  $b$ , exponent  $e$ , modulus  $n$ 
output:  $b^e \pmod{n}$ 
 $X \leftarrow 1$ ;
for  $i \leftarrow \text{bitlen}(e)$  downto 0 do
     $X \leftarrow \text{multiply}(X, X)$ ;
    if  $e_i = 1$  then  $X \leftarrow \text{multiply}(X, b)$ ;
end
end
return  $X$ ;
```

---

# Attacks

## Page Table Based Attacks

- Basic page-fault attack
  - AEX on fault  $\Rightarrow$  base address of faulting page revealed
  - Cause intentional page faults  $\Rightarrow$  obtain page level trace
  - Monitor a specific page



```
void CountLogin( GENDER s ) {
    if ( s == MALE ) {
        gMaleCount ++;
    } else {
        gFemaleCount ++;
    }
}
```

```
char* WelcomeMessage( GENDER s ) {
    char *mesg;

    // GENDER is an enum of MALE and FEMALE
    if ( s == MALE ) {
        mesg = WelcomeMessageForMale();
    } else { // FEMALE
        mesg = WelcomeMessageForFemale();
    }
    return mesg;
}
```



# Attacks

---

## *Page Table Based Attacks*

- Problem: page-faults cause a high overhead => attack detection
  - Solution: Use page-table access and dirty bits
- Problem: TLB caches PTEs and flags are not updated in PT
  - Solution: Flush TLB using an IPI
- Problem: Lots of IPIs
  - Solution: Time difference between two repeatedly accessed pages  $\Rightarrow$  one interrupt
- Problem: Still interrupts!
  - Solution: TLB is shared between two hyperthreads on same core  $\Rightarrow$  invalidate TLB entries  $\Rightarrow$  no IPIs

# Attacks

---

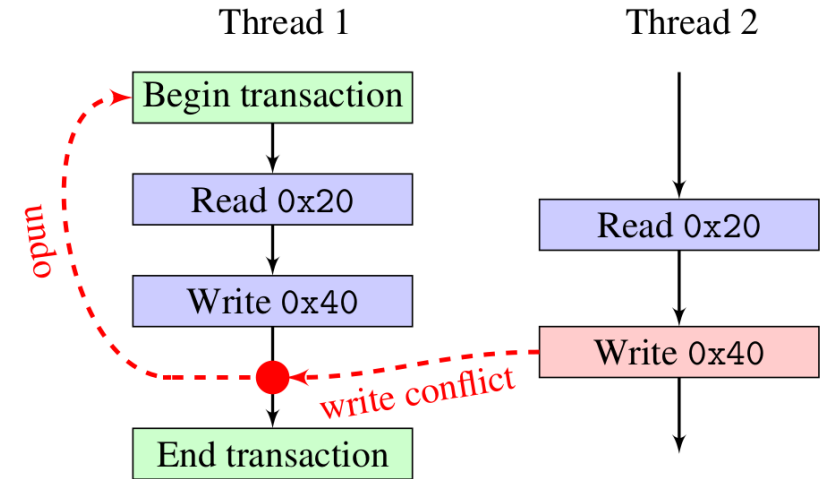
## *Other Attacks*

- Flush+Reload
  - Flush a memory line and wait for victim to access it
- Flush+Flush
  - Flush a cache line using clflush instruction
- DRAMA
  - Caching disabled
  - Allocate two memory lines with different virtual addresses but same physical address
  - Regularly access one of the memory lines
- Cache-DRAM
  - Prime+Probe + DRAMA

# Defenses

## *Defense against cache-based attacks*

- Basic approach: Pin code/data to cache
- Using Intel TSX
  - Introduced with Haswell
  - Critical sections management
  - Monitoring serialization
  - Restricted Transactional Memory
    - New Instruction set interface
- Cloak
  - Preload code/data
  - Execute algorithm inside transaction

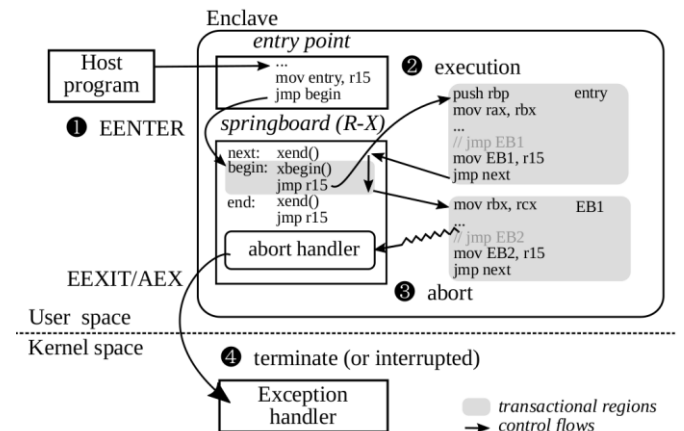



```
xbegin();  
preload all sensitive data/code;  
run algorithm;  
xend();
```

# Defenses

## *Defense against page table based attacks*

- T-SGX: run code inside transaction => page fault => attack detection => program termination
- Problem: Frequent timer interrupt
  - Program take too long => exception => never terminate
- Problem: TSX buffers all memory changes inside cache
  - single transaction, lots of accesses => cache conflict => transaction abort
- Break program into tiny blocks
  - Problem: transaction setup has cost => performance
- Merge continuous blocks
- Problem: boundaries leak
  - Solution: springboard



# Defenses

---

## *Other defense methods*

- Déjà vu
  - Reference clock thread inside a transaction
- Shinde et al.
  - Deterministic page access profile
  - Introducing fake acceses
- SGX-Shield
  - ASLR
- ZeroTrace
  - Use of ORAM + SGX
- Dr. SGX
  - Data location randomization

# Analysis

---

- Cloak
  - Makes cache-pinning possible to some extent
  - Execution time can leak
  - Aborts do not cancel concurrent memory accesses
  - Execution behavior and branch prediction uncertainties
- T-SGX
  - Attacks based on monitoring flags are possible
- Déjà vu
  - Page monitoring attacks still effective
  - Only works on AEX based attacks
- Bottom line
  - Transaction based defenses usually come with high overload and low utilization, and need isolation

# Analysis

- Other methods
  - Attack detection methods  $\Rightarrow$  Unreliable
  - Shuffling memory  $\Rightarrow$  expensive
  - ORAM (make addresses input-independent)  $\Rightarrow$  expensive
- The perfect solution?
  - Remove all branches and conditional code
  - Pin data/code to cache and TLB
  - CAT+SGX
  - Non-shared non-cached secure memory element
- Most of defenses either work on page table based attacks or cache-based attacks
- ORAM performance optimization
- Page monitoring attacks

Attack/Defense	P+P	Basic PF	F+R	F+F	DRAMA	PM
Cloak	✓	✗	✓	✓	✗	✗
T-SGX	✗	✓	✗	✗	✓	✗
Deja Vu	✗	✓	✗	✗	✓	✗
Shinde et al.	✗	✓	✗	✗	✓	✓
SGX-Shield	✗	✓	✗	✗	✓	✗
ZeroTrace	✓	✓	✓	✓	✓	✓
DR. SGX	✓	✗	✓	✓	✗	✗

# Conclusion

---

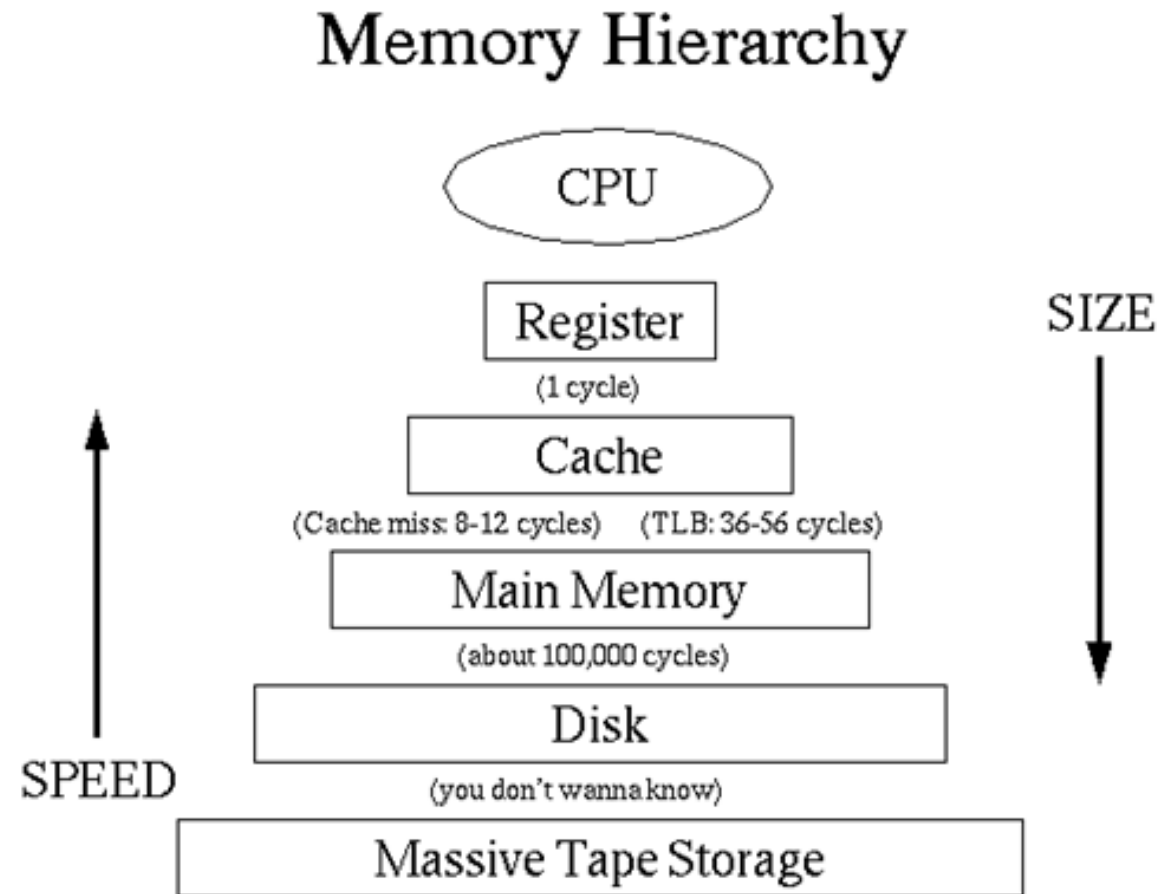
- We studied multiple attacks and defenses
- Still no robust defense
- Attack are emerging
- Open to research



# References

- 
- [1] Ferdinand Brasser, Srdjan Capkun, Alexandra Dmitrienko, Tommaso Frassetto, Kari Kostinen, Urs Müller, and Ahmad-Reza Sadeghi. 2017. DR. SGX: Hardening SGX Enclaves against Cache Attacks with Data Location Randomization. arXiv preprint arXiv:1709.09917 (2017).
  - [2] Ferdinand Brasser, Urs Müller, Alexandra Dmitrienko, Kari Kostinen, Srdjan Capkun, and AhmadReza Sadeghi. 2017. Software grand exposure: SGX cache attacks are practical. arXiv preprint arXiv:1702.07521 (2017), 33.
  - [3] Sanchuan Chen, Xiaokuan Zhang, Michael K Reiter, and Yinqian Zhang. 2017. Detecting privileged side-channel attacks in shielded execution with D'éj'a Vu. In Proceedings of the 2017 ACM on Asia Conference on Computer and Communications Security. ACM, 7–18.
  - [4] Oded Goldreich. 1987. Towards a theory of software protection and simulation by oblivious RAMs. In Proceedings of the nineteenth annual ACM symposium on Theory of computing. ACM, 182–194.
  - [5] Oded Goldreich and Rafail Ostrovsky. 1996. Software protection and simulation on oblivious RAMs. Journal of the ACM (JACM) 43, 3 (1996), 431–473.
  - [6] Daniel Gruss, Julian Lettner, Felix Schuster, Olya Ohrimenko, Istvan Haller, and Manuel Costa. 2017. Strong and efficient cache side-channel protection using hardware transactional memory. In USENIX Security Symposium.
  - [7] Daniel Gruss, Clémentine Maurice, Klaus Wagner, and Stefan Mangard. 2016. Flush+ Flush: a fast and stealthy cache attack. In International Conference on Detection of Intrusions and Malware, and Vulnerability Assessment. Springer, 279–299.
  - [8] Mehmet Kayaalp, Nael Abu-Ghazaleh, Dmitry Ponomarev, and Aamer Jaleel. 2016. A high-resolution side-channel attack on last-level cache. In Proceedings of the 53rd Annual Design Automation Conference. ACM, 72.
  - [9] Rafail Ostrovsky. 1990. Efficient computation on oblivious RAMs. In Proceedings of the twenty-second annual ACM symposium on Theory of computing. ACM, 514–523.
  - [10] Dag Arne Osvik, Adi Shamir, and Eran Tromer. 2006. Cache attacks and countermeasures: the case of AES. In Cryptographers' Track at the RSA Conference. Springer, 1–20.
  - [11] Peter Pessl, Daniel Gruss, Clémentine Maurice, Michael Schwarz, and Stefan Mangard. 2016. DRAMA: Exploiting DRAM Addressing for Cross-CPU Attacks.. In USENIX Security Symposium. 565–581.
  - [12] Sajin Sasy, Sergey Gorbunov, and Christopher Fletcher. 2017. ZeroTrace: Oblivious memory primitives from Intel SGX. In Symposium on Network and Distributed System Security (NDSS).
  - [13] Jaebaek Seo, Byoungyoung Lee, Seongmin Kim, Ming-Wei Shih, Insik Shin, Dongsu Han, and Taesoo Kim. 2017. SGX-Shield: Enabling address space layout randomization for SGX programs. In Proceedings of the 2017 Annual Network and Distributed System Security Symposium (NDSS), San Diego, CA.
  - [14] Ming-Wei Shih, Sangho Lee, Taesoo Kim, and Marcus Peinado. 2017. T-SGX: Eradicating controlledchannel attacks against enclave programs. In Proceedings of the 2017 Annual Network and Distributed System Security Symposium (NDSS), San Diego, CA.
  - [15] Shweta Shinde, Zheng Leong Chua, Viswesh Narayanan, and Prateek Saxena. 2015. Preventing your faults from telling your secrets: Defenses against pigeonhole attacks. arXiv preprint arXiv:1506.04832 (2015).
  - [16] Eran Tromer, Dag Arne Osvik, and Adi Shamir. 2010. Efficient cache attacks on AES, and countermeasures. Journal of Cryptology 23, 1 (2010), 37–71.
  - [17] Jo Van Bulck, Nico Weichbrodt, Rüdiger Kapitza, Frank Piessens, and Raoul Strackx. 2017. Telling your secrets without page faults: Stealthy page table-based attacks on enclaved execution. In Proceedings of the 26th USENIX Security Symposium. USENIX Association.
  - [18] Wenhao Wang, Guoxing Chen, Xiaorui Pan, Yinqian Zhang, Xiaofeng Wang, Vincent Bindschaedler, Haixu Tang, and Carl A. Gunter. 2017. Leaky Cauldron on the Dark Land: Understanding Memory Side-Channel Hazards in SGX. 2421–2434.
  - [19] Yuval Yarom and Katrina Falkner. 2014. FLUSH+ RELOAD: A High Resolution, Low Noise, L3 Cache Side-Channel Attack.. In USENIX Security Symposium

# Introduction



# Defenses

---

- SGX-Shield
  - Address space layout randomization
- ORAM+SGX
- Data shuffling

# Introduction

---

## *Intel TSX*

- Introduced with Haswell
- Critical sections management
- Monitoring serialization
- Restricted Transactional Memory
  - New Instruction set interface

# Attacks

---

## *Points of attack*

- TLB: shared between enclave and non-enclave programs  $\Rightarrow$  with hyperthreading enabled it creates side-channels
- Page faults: visible to OS
- TLB flushes: on context switch
- Page table entry flags
- Enclave memory address beginning and offset are known
- Enclave exits (AEX)

# Analysis

---

- SGX-Shield
  - No live randomization  $\Rightarrow$  observe and monitor random patterns
- Shinde et al.
  - Cache and TLB based attacks possible